

1

**Outline**Volumes and Networking

- Create and manage volumes
- Sharing Data using volumes
- Volumes vs Mounts
- The Docker network topology
- Understanding the default network
- Bridge vs Overlay networks

M.Romdhani, 2020

2

2

## Create and manage volumes

3

### Why Docker volumes ?

Volumes and Networking

- Container filesystems are ephemeral
- Docker volumes allow data to persist beyond the life of a container
- Containers don't 'own' volumes, but can reference them for use
- Volumes can be shared by all containers running on the Docker host, and may contain data from the host

M.Romdhani, 2020

4

4

## Volumes are special directories in a container

Volumes and Networking

### ■ Volumes can be declared in two different ways:

1. Within a Dockerfile, with a VOLUME instruction.  
■ `VOLUME /uploads`
2. On the command-line, with the -v flag for docker run.  
■ `$ docker run -d -v /uploads myapp`

### ■ In both cases, `/uploads` (inside the container) will be a volume.

M.Romdhani, 2020

5

5

## Types of Docker Volumes

Volumes and Networking

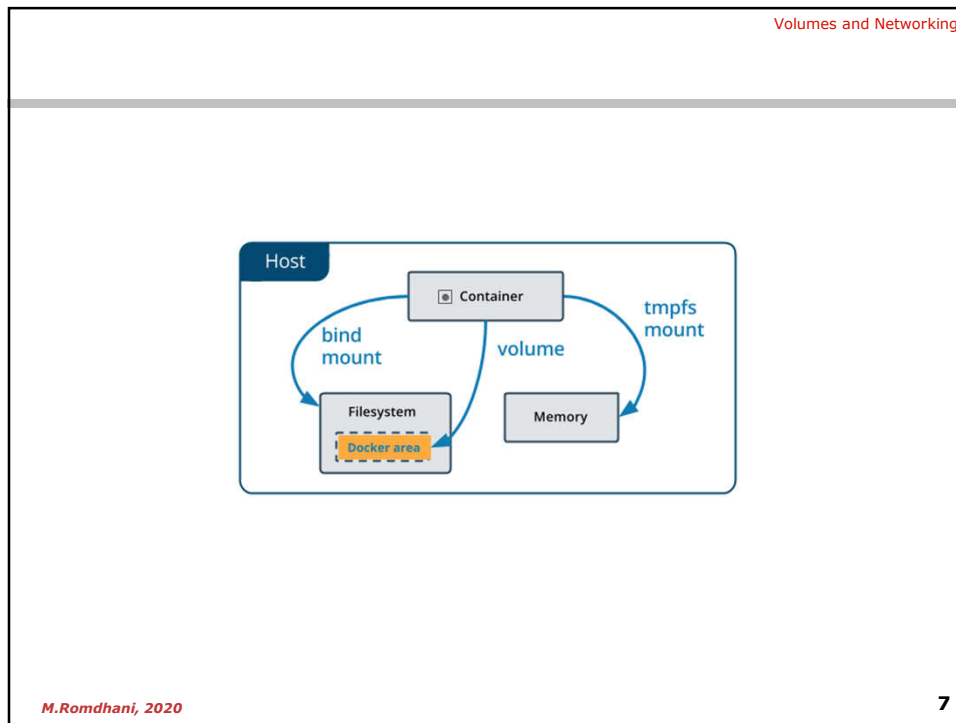
### ■ There are three types of volumes:

1. **Named volumes** : independent volume entities, created and managed independently of containers
2. **Container volume**: volumes created in conjunction with a specific container
3. **Host directory/file bind mount**: not strictly a volume, but a means of sharing data with a container from the host

M.Romdhani, 2020

6

6



7

Volumes and Networking

## Named Volumes

- **Named volumes are created independently of containers with the docker volume create command**
- **Named volumes can be created with the default driver, or with a third-party plugin driver**
- **Format of the docker volume create sub-command is:**
  - `docker volume create [options]`
- **The config options available for docker volume create:**

Client Options	Description
<code>-d, --driver=local</code>	Driver to use when creating the named volume
<code>--name=""</code>	Name to be applied to the volume
<code>-o, --opt=map[]</code>	Driver specific options to apply to the volume

8

M.Romdhani, 2020

8

## Using Named Volumes

- **Named volumes are used by containers when referenced as an argument to the `-v`, `--volume` config option of the Docker CLI:**
  - `$ docker run -v archive:/backup busybox tar cvf \`  
`> /backup/archive-2015-03-11.tar /data`
- **The named volume must be specified *without a preceding '/'*, which provides another different meaning**
- **If the named volume referenced in a Docker CLI command doesn't exist, it will be created with the container**

## Container Volumes

- **Docker volumes can be created dynamically at container runtime**
- **The volume is created on the host and mounted into the container at the specified mount point**
- **Control parameters can be applied to characterize the volume's use (e.g. `ro|rw`)**
- **Docker volumes use pluggable drivers, and third-party drivers are available**

Client Option	Description
<code>-v</code> , <code>--volume</code>	Specifies a volume to mount into container filesystem
<code>--volume-driver</code>	Driver to use when creating and operating the volume

## Host Files and Directories

- In addition to volumes, files & directories from the host can be mounted into containers
- The `-v`, `--volume` config option is used for this purpose:
  - `$ docker run -d -v /data:/var/lib/mysql mariadb`
- If container location exists, the host file or directory is mounted over the top of the container location
- The following control parameters can be applied as part of the `-v`, `--volume` config option:

Control Parameters	Description
<code>rw ro</code>	Specifies whether mount is read-write or read-only
<code>z Z</code>	For SELinux labels, specifies if mount is private or not
<code>[r]shared [r]slave [r]private</code>	Sets propagation properties of bind mount

## docker volume rm & ls

- The `docker volume rm` command removes volumes
  - Volumes cannot be removed if they are in use
  - The format of the docker volume rm sub-command is:
    - `docker volume rm volume [volume ...]`
- The `docker volume ls` command is for listing volumes
  - The format of the docker volume ls sub-command is:
    - `docker volume ls [options]`
  - Config options available for docker volume ls:

Client Options	Description
<code>-f, --filter=[]</code>	Filter output based on set criteria (dangling=true)
<code>-q, --quiet=false</code>	Only display volume names in the output

## docker volume inspect

- The docker volume inspect sub-command provides detailed information for the specified volume(s)
- The format of the docker volume inspect sub-command is:
  - `docker volume inspect [option] volume [volume ...]`
- The single config option available for docker volume inspect is:

Client Options	Description
<code>-f, --format=""</code>	Golang text/template to apply to format the output

**Share volumes across  
containers**

## Volumes bypass the copy-on-write system

### ■ Volumes act as passthroughs to the host filesystem

- The I/O performance on a volume is **exactly** the same as I/O performance on the Docker host.
- When you docker commit, the content of volumes is **not brought into the resulting image**.
- If a RUN instruction in a Dockerfile changes the content of a volume, those changes are not recorded neither.
- If a container is started with the **--read-only** flag, the volume will still be writable (unless the volume is a read-only volume).

## Volumes can be shared across containers

- You can start a container with **exactly the same volumes** as another one.
- The new container will have the same volumes, in the same directories.
- They will contain exactly the same thing, and remain in sync.
- Under the hood, they are actually the same directories on the host anyway.
- This is done using the **--volumes-from** flag for **docker run**.
- We will see an example in the following slides.



## Sharing app server logs with another container

Volumes and Networking

- Let's start a Tomcat container:

- `$ docker run --name webapp -d -p 8080:8080 -v /usr/local/tomcat/logs tomcat`

- Now, start an alpine container accessing the same volume:

- `$ docker run --volumes-from webapp alpine sh -c "tail -f /usr/local/tomcat/logs/*"`

- Then, from another window, send requests to our Tomcat container:

- `$ curl localhost:8080`

M.Romdhani, 2020

17

17

## Sharing Data using Named volumes

Volumes and Networking

- Volumes can be created without a container, then used in multiple containers.

- Let's create a couple of volumes directly.

- `$ docker volume create webapps`

- `$ docker volume create logs`

- Volumes are not anchored to a specific path.

M.Romdhani, 2020

18

18

## Using our named volumes

- Volumes are used with the `-v` option.
- When a host path does not contain a `/`, it is considered a volume name.
- Let's start a web server using the two previous volumes.
 

```
$ docker run -d -p 1234:8080 \
    -v logs:/usr/local/tomcat/logs \
    -v webapps:/usr/local/tomcat/webapps \
    tomcat
```
- Check that it's running correctly:
 

```
$ curl localhost:1234
```

  - ... (Tomcat tells us how happy it is to be up and running) ...

M.Romdhani, 2020

19

19

## Using a volume in another container

- We will make changes to the volume from another container.
- In this example, we will run a text editor in the other container.
- Let's start another container using the webapps volume.
 

```
$ docker run -v webapps:/webapps -w /webapps -ti alpine vi
ROOT/index.jsp
```
- Change the page, save, exit.
- Then run `curl localhost:1234` again to see your changes

M.Romdhani, 2020

20

20

## Migrating data with --volumes-from

- The `--volumes-from` option tells Docker to re-use all the volumes of an existing container.
  - Scenario: migrating from Redis 2.8 to Redis 3.0.
  - We have a container (myredis) running Redis 2.8.
  - Stop the myredis container.
  - Start a new container, using the Redis 3.0 image, and the `--volumes-from` option.
  - The new container will inherit the data of the old one.
  - Newer containers can use `--volumes-from` too.
  - Doesn't work across servers, so not usable in clusters (Swarm, Kubernetes).

## Data migration in practice

- Let's create a Redis container.
 

```
$ docker run -d --name redis28 redis:2.8
```
- Connect to the Redis container and set some data.
 

```
$ docker run -ti --link redis28:redis busybox telnet redis 6379
```
- Issue the following commands:
 

```
SET counter 42
INFO server
SAVE
QUIT
```

Volumes and Networking

## Data migration in practice

---

- **Upgrading Redis**
  - Stop the Redis container.  
`$ docker stop redis28`
  - Start the new Redis container.  
`$ docker run -d --name redis30 --volumes-from redis28 redis:3.0`
- **Testing the new Redis**
  - Connect to the Redis container and see our data.  
`docker run -ti --link redis30:redis busybox telnet redis 6379`
- **Issue a few commands.**

```
GET counter
INFO server
QUIT
```

M.Romdhani, 2020
23

23

Volumes and Networking

## Using custom "bind-mounts"

---

- **In some cases, you want a specific directory on the host to be mapped inside the container:**
  - You want to manage storage and snapshots yourself.
  - You have a separate disk with better performance (SSD) or resiliency (EBS) than the system disk, and you want to put important data on that disk.
  - You want to share your source directory between your host (where the source gets edited) and the container (where it is compiled or executed).
- **Wait, we already met the last use-case in our example development workflow! Nice.**
  - `$ docker run -d -v /path/on/the/host:/path/in/container image`

M.Romdhani, 2020
24

24

## Volumes vs Mounts

25

## Volumes vs. Mounts

Volumes and Networking

- Since Docker 17.06, a new options is available: `--mount`.
- It offers a new, richer syntax to manipulate data in containers.
- It makes an explicit difference between:
  - volumes (identified with a unique name, managed by a storage plugin),
  - bind mounts (identified with a host path, not managed).
- The former `-v` / `--volume` option is still usable.

M.Romdhani, 2020

26

26

## --mount syntax

### ■ Binding a host path to a container path:

```
$ docker run \  
--mount type=bind,source=/path/on/host,target=/path/in/container alpine
```

### ■ Mounting a volume to a container path:

```
$ docker run \  
--mount source=myvolume,target=/path/in/container alpine
```

### ■ Mounting a tmpfs (in-memory, for temporary files):

```
$ docker run \  
--mount type=tmpfs,destination=/path/in/container,tmpfs-size=1000000  
alpine
```

## The Docker network topology

## Container Network Model (CNM)

- The CNM was introduced in Engine 1.9.0
  - Up until Docker Engine 1.9 (2015), native Docker networking was confined to a single host
- Containers can share an overlay network, a local bridged network, the host's stack, or a third-party plugin network
- The CNM adds the notion of a network, and a new top-level command to manipulate and see those networks: `docker network`.

```
$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
70dfd633ba3b	bridge	bridge	local
9c7a9895e729	host	host	local
585f508793d7	none	null	local

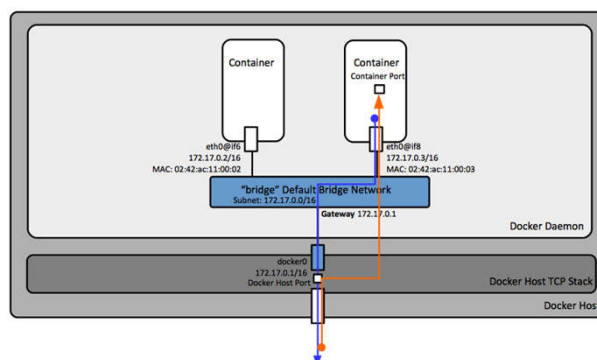
M.Romdhani, 2020

29

29

## Container networking

- Docker will allocate IP addresses to the containers connected to a network.
  - Containers can be connected to multiple networks.
  - Containers can be given per-network names and aliases.
  - The names and aliases can be resolved via an embedded DNS server.



M.Romdhani, 2020

30

30

## Network implementation details

- A network is managed by a driver.
- The built-in drivers include:
  - `bridge` (default)
  - `none`
  - `host`
  - `macvlan`
- A multi-host driver, `overlay`, is available out of the box (for Swarm clusters).
- More drivers can be provided by plugins (OVS, VLAN...)

## Creating a network

- Let's create a network called dev.
 

```
$ docker network create dev
4c1ff84d6d3f1733d3e239cac276f425a9d5228a4355d54878293a889ba
```
- The network is now visible with the `network ls` command:

```
$ docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
4c1ff84d6d3f        dev                 bridge              local
70dfd633ba3b        bridge             bridge              local
9c7a9895e729        host               host                local
585f508793d7        none              null                local
```



## Placing containers on a network

- We will create a named container on this network.

- It will be reachable with its name, **es**

```
$ docker run -d --name es --net dev elasticsearch:2
8abb80e229ce8926c7223beb69699f5f34d6f1d438bfc5683e798046863
```

## Communication between containers

- Now, create another container on this network.

```
$ docker run -it --net dev alpine sh
root@0ecccdfa45ef:/#
```

- From this new container, we can resolve and ping the other one, using its assigned name:

```
/ # ping es
PING es (172.18.0.2) 56(84) bytes of data.
64 bytes from es.dev (172.18.0.2): icmp_seq=1 ttl=64 time=0.221 ms
64 bytes from es.dev (172.18.0.2): icmp_seq=2 ttl=64 time=0.114 ms
64 bytes from es.dev (172.18.0.2): icmp_seq=3 ttl=64 time=0.114 ms
^C
--- es ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2000ms
rtt min/avg/max/mdev = 0.114/0.149/0.221/0.052 ms
root@0ecccdfa45ef:/#
```

## Understanding the default network

35

### A simple, static web server

Volumes and Networking

- Run the Docker Hub image nginx, which contains a basic web server:

```
$ docker run -d -P nginx
```

```
66b1ce719198711292f84a7b68c3876cf9f67015e752b94e189d35a204e
```

- Docker will download the image from the Docker Hub.

- **-d** tells Docker to run the image in the background (Detached mode).
- **-P** tells Docker to make this service reachable from other computers. (-P is the short version of --publish-all.)

But, how do we connect to our web server now?

M.Romdhani, 2020

36

36

## Finding our web server port

- We will use docker ps:

```
$ docker ps
CONTAINER ID  IMAGE  ...  PORTS  ...
e40ffb406c9e  nginx  ...  0.0.0.0:32768->80/tcp  ...
```

- The web server is running on port 80 inside the container.
- This port is mapped to port 32768 on our Docker host.

- We can also use docker port:

```
$ docker port e40ffb406c9e
80/tcp -> 0.0.0.0:32768
```

## Connecting to our web server (CLI)

- You can also use curl directly from the Docker host.

- Make sure to use the right port number if it is different from the example below:

```
$ curl localhost:32768
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
...
```

## How does Docker know which port to map?

- There is metadata in the image telling "this image has something on port 80".
- We can see that metadata with docker inspect:  

```
$ docker inspect --format '{{.Config.ExposedPorts}}' nginx map[80/tcp:{}]
```
- This metadata was set in the Dockerfile, with the **EXPOSE** keyword.
- We can see that with docker history:

```
$ docker history nginx
IMAGE          CREATED          CREATED BY          CMD
7f70b30f2cc6   11 days ago     /bin/sh -c #(nop)   CMD ["nginx" "-g" "...
<missing>      11 days ago     /bin/sh -c #(nop)   STOPSIGNAL [SIGTERM]
<missing>      11 days ago     /bin/sh -c #(nop)   EXPOSE 80/tcp
```

## Manual allocation of port numbers

- If you want to set port numbers yourself, no problem:  

```
$ docker run -d -p 80:80 nginx
$ docker run -d -p 8000:80 nginx
$ docker run -d -p 8080:80 -p 8888:80 nginx
```
- We are running three NGINX web servers.
  - The first one is exposed on port 80.
  - The second one is exposed on port 8000.
  - The third one is exposed on ports 8080 and 8888.

## Finding the container's IP address

- We can use the `docker inspect` command to find the IP address of the container.

```
$ docker inspect --format '{{ .NetworkSettings.IPAddress }}'
<yourContainerID>
172.17.0.3
```

- `docker inspect` is an advanced command, that can retrieve a ton of information about our containers.
- Here, we provide it with a format string to extract exactly the private IP address of the container.

## Looking at the network setup in the container

- We can look at the list of network interfaces with `ifconfig`, `ip a`, or `ip l`:

```
/ # ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
18: eth0@if19: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
    link/ether 02:42:ac:11:00:02 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.2/16 brd 172.17.255.255 scope global eth0
        valid_lft forever preferred_lft forever
20: eth1@if21: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
    link/ether 02:42:ac:14:00:04 brd ff:ff:ff:ff:ff:ff
    inet 172.20.0.4/16 brd 172.20.255.255 scope global eth1
        valid_lft forever preferred_lft forever
/ #
```

- Each network connection is materialized with a virtual network interface.
- As we can see, we can be connected to multiple networks at the same time.

## Bridge vs Overlay networks


43

### Bridge Networking

Volumes and Networking

- At start-up, the Docker daemon creates three default networks, called 'bridge', 'none', and 'Container'

- Bridge
  - By default, any container invoked will be connected to this bridge, with its network stack contained within its own NET namespace, but with its primary network interface connected to the bridge
- None
  - With `--net=none`, the container is created with just a loopback interface to allow container-specific services to communicate
- Host
  - The container shares the same NET namespace and network stack (interfaces, ports etc.) as the host
- Container
  - The container shares the network stack of another specified container



M.Romdhani, 2020

44

44

## Creating Bridge Networks

Volumes and Networking

- To create a local user-defined bridge network:

```
$ docker network create -d bridge local_bridge
```

- The `--link` config option provides a private alias:

```
$ docker run -d --name provider --net local_bridge gcr.io/google-containers/pause
a21d9881993463276834e9e962da43657cb898b764c60a25375b5341b193c627
$ docker run -it --name consumer --link provider:giver --net local_bridge busybox sh
/ # ping -q -c 1 provider
PING provider (172.22.0.2): 56 data bytes

--- provider ping statistics ---
1 packets transmitted, 1 packets received, 0% packet loss
round-trip min/avg/max = 0.110/0.110/0.110 ms
/ # ping -q -c 1 giver
PING giver (172.22.0.2): 56 data bytes

--- giver ping statistics ---
1 packets transmitted, 1 packets received, 0% packet loss
round-trip min/avg/max = 0.166/0.166/0.166 ms
```

M.Romdhani, 2020


45

45

## Overlay networks for multi host networks

Volumes and Networking

- Docker's native networking now allows for the creation of overlay networks which span multiple Docker hosts, allowing containers running on different Docker hosts to communicate as if they were co-hosted
  - Overlay networks are VXLANs connecting different Docker hosts, and require a key/value store to hold state



M.Romdhani, 2020

46

46

## Creating Overlay networks

### ■ Main Steps:

- Enable Swarm Mode (`docker swarm init` then `docker swarm join` on other nodes)
- `docker network create mynet --driver overlay`
- `docker service create --network mynet myimage`